



Tribhuvan University
Faculty of Humanities and Social Science
Swastik college, Bhaktapur



Distributed System
Case Study on
The Global Name Service
Submitted By:
Safalta Shrestha (113102089)
BCA 6th Semester

Submitted To:
Sagar Rana Magar
Lecturer, BCA Department
Swastik college, Bhaktapur
January 2025

Table Of Content

Executive Summary.....	III
Chapter-1 Introduction and Background.....	1
1.1 Introduction.....	1
1.2 Background.....	1
1.3 Problem Statement.....	2
1.4 Objectives.....	2
Chapter-2 Architecture of GNS.....	3
2.1 Directory Tree and Value Tree.....	3
2.2 Directory Identifier(DI).....	3
2.3 Name of Entry.....	3
2.4 Multiple Name Servers.....	3
2.5 Replication.....	4
2.6 Caching.....	4
Chapter-3 Working of GNS.....	5
Chapter-4 Application and Importance of GNS.....	7
4.1 Application.....	7
4.2 Importance.....	7
Chapter-5 Limitation of GNS.....	9
Chapter-6 Real-World Implementation of GNS.....	10
Chapter-7 Conclusion.....	11
References.....	12

Executive Summary

The Global Name Service (GNS) is a distributed and scalable naming system designed to map human-readable names to network resources, including IP addresses, service endpoints, email identifiers, and other attributes. Unlike traditional naming systems, GNS provides a hierarchical directory structure with separate value trees to store resource information. Each directory is uniquely identified by a Directory Identifier (DI), enabling safe renaming, merging, and movement of directories without breaking references.

GNS leverages replication and caching to enhance reliability, availability, and performance across distributed environments. Its design supports dynamic, large-scale systems, making it suitable for cloud computing, Internet of Things (IoT) networks, and decentralized platforms. By separating names from physical locations and providing a flexible global namespace, GNS ensures efficient resource discovery, secure name resolution, and scalable management in modern distributed systems.

Chapter-1 Introduction and Background

1.1 Introduction

With the rapid growth of computer networks and distributed systems, the process of identifying, locating, and accessing network resources has become increasingly complex. In the early days, resources were accessed using numerical network addresses, which were difficult for users to remember and inefficient to manage as networks expanded. As the number of devices, services, and users increased globally, this approach proved unsuitable for large-scale distributed environments.

To address these challenges, naming systems were introduced to provide human-readable names that could be mapped to underlying network resources. The Global Name Service (GNS) was developed as a distributed naming framework designed to support scalable and flexible name resolution. GNS provides a global namespace that allows users and applications to locate resources without needing to know their physical location or network address.

Unlike traditional naming approaches that focus on a limited set of resource mappings, GNS supports a wide range of resource attributes, including IP addresses, service endpoints, email identifiers, and security-related information. By separating naming from physical location and using a hierarchical directory structure, GNS improves usability, scalability, and manageability. As a result, GNS plays a vital role in modern distributed systems, cloud computing environments, and emerging decentralized networks.

1.2 Background

In the early development of computer networks, naming and addressing information was managed using centralized mechanisms, such as static host tables. These approaches were effective only for small networks and quickly became inefficient as the number of connected systems increased. Centralized management led to scalability issues, frequent inconsistencies, and administrative difficulties, highlighting the need for more robust naming solutions.

To overcome these limitations, distributed naming systems were introduced, enabling network identifiers to be managed across multiple servers in a hierarchical manner. Building on these foundational ideas, the Global Name Service (GNS) was developed to provide a more flexible and extensible naming architecture suitable for large-scale distributed environments.

GNS introduces the concept of separating directory structures from value storage using directory trees and value trees. Each directory is uniquely identified by a Directory Identifier (DI), allowing directories to be renamed, moved, or merged without breaking existing references. This design enhances consistency, scalability, and ease of management. By supporting replication, caching, and distribution across multiple servers, GNS ensures high availability and reliable name resolution, making it well-suited for modern applications such as cloud services, Internet of Things (IoT), and decentralized systems.

1.3 Problem Statement

Modern distributed systems require a naming mechanism that is scalable, flexible, and capable of handling dynamic resources. Traditional naming approaches often lack support for decentralized control, secure management, and diverse resource attributes. With the rapid growth of cloud computing, Internet of Things (IoT), and decentralized systems, these limitations become more pronounced. The Global Name Service (GNS) addresses these challenges by providing a distributed, extensible, and reliable global naming framework suitable for modern networking environments.

1.4 Objectives

The primary objectives of this study are:

- To understand the concept and architecture of the Global Name Service (GNS)
- To study the working mechanism of GNS and its components
- To analyze the advantages and limitations of GNS
- To compare GNS with the traditional naming systems such as the Domain Name System (DNS)
- To explore real-world applications of GNS in modern distributed systems

Chapter-2 Architecture of GNS

The Global Name Service(GNS) is designed to provide scalable, distributed and reliable name resolution mechanism. Its architecture is based on hierarchical naming, distribution across multiple servers, replication and caching.

2.1 Directory Tree and Value Tree

- GNS organizes names in a directory tree similar to a file system.
- Each directory may contain sub-directories and entries that reference value trees.
- The value tree is located at the leaf node and stores actual attributes such as IP address, password, email ID and service endpoint.
- Directory does not store values directly, instead they maintain references to the corresponding value trees allowing flexibility and scalability.

2.2 Directory Identifier(DI)

- Each directory in GNS is assigned a Directory Identifier (DI).
- DI is a unique identifier, that is independent of the directory name.
- Advantages of using DI include: directories can be renamed without affecting existing references, safe merging of directory trees is possible and, improves consistency and scalability.

2.3 Name of an Entry

A GNS name consists of two parts directory name and value name. This structure allows precise identification and retrieval of the required resource within the directory hierarchy.

2.4 Multiple Name Servers

- The global directory tree is partitioned and distributed across multiple name servers.
- No single server stores the entire name-space.

- Each server is responsible for managing and resolving a specific portion of the directory tree, enabling load distribution and scalability.

2.5 Replication

- Each directory partition is replicated across several servers.
- Replication provides fault tolerance, high availability , reliability in case of server failures

2.6 Caching

- Frequently accessed directory entries and values are cached to improve performance.
- Cached data may become temporarily inconsistent, which is acceptable in GNS
- This trade-off improves response time and enables faster name solution




Figure 2.1 Architecture of GNS

Chapter-3 Working of GNS

The Global Name System (GNS) establishes the connection between human-readable names and network resource identifiers such as IP addresses by storing the IP address as an attribute in the leaf nodes of a hierarchical directory tree. GNS operates in two main phases:

1. Name Creation (Publishing Phase)

This phase involves creating and storing the mapping between a human-readable name and the corresponding network address.

2. Name Finding (Resolution Phase)

This phase retrieves the network address when a user queries a human-readable name.

The following steps explain how a name is created in GNS and how the corresponding result is retrieved by the user.

Step 1: Create a Directory in the GNS Tree

An administrator creates a directory node within the GNS hierarchy. Each directory is uniquely identified by a Directory Identifier(DI) , independent of its name.

Step 2: Assign a Name

Inside the directory, a local name like a file name in a folder is created. This represents an individual resource or service entry under that directory.

Step 3: Store IP Address as a Value (Attribute)

The IP address is stored as an attribute within a value at the leaf node of the directory. The leaf contains structured information such as IP address,password, email ID or other resource attributes .

Step 4: Directory Entry Points to the Value Tree

The directory entry itself does not store the IP directly. Instead, it holds a reference to the value tree where the actual IP address and other attributes reside. This separation improves flexibility and scalability.

Step 5: Store Directory on Authoritative Server

Each directory and its entries are managed by an authoritative GNS server. Directory data is replicated and cached across multiple servers to ensure high availability and fault tolerance.

Step 6: Name Resolution Creates the Connection at Runtime

- When a client looks up the name:
 - The query starts at the root or working directory.
 - It follows directory references step by step through the hierarchical tree.
 - It navigates through the directory nodes to reach the leaf node.
 - The IP address attribute is read from the value tree at the leaf.
 - The resolved IP address is returned to the client for use.

Chapter-4 Application and Importance of GNS

4.1 Application

Some applications include:

1. Global Name Resolution: GNS enables the resolution of human-readable names to various network resources such as IP addresses, service endpoints and identifiers, simplifying resource access in distributed systems.

2. Distributed System and Cloud Services: In large distributed environments and cloud platforms, GNS provides a unified namespace for locating resources such as servers, databases, storage services, and microservices across geographically dispersed data centers.

3. Service Discovery: Applications use GNS to dynamically discover network services and endpoints. This is essential in environments with frequent service updates, such as container orchestration platforms (e.g., Kubernetes).

4. Communication and Messaging Systems: GNS supports the identification and location of communication services by mapping human-readable names to appropriate service endpoints, enabling reliable message routing in distributed environments.

5. Network Management: Network devices like routers, switches, and firewalls use GNS for naming and managing network elements, enhancing ease of configuration and troubleshooting.

6. Internet of Things(IoT): IoT ecosystems, which involve billions of connected devices, use GNS for naming and addressing devices to enable efficient communication and management.

7. Security and Access Control: GNS supports naming schemes that help implement security policies, such as access control lists and certificates, by providing verifiable resource identifiers.

4.2 Importance

The Global Name Service is fundamental to the operation and growth of the Internet and modern networks due to the following reasons:

1. Scalability: GNS's distributed architecture supports billions of devices and services, scaling efficiently as the Internet expands.

2. Human-Friendly Naming: By allowing human-readable names, GNS simplifies interaction with distributed resources and services reducing the complexity of using numerical identifiers.

3. Reliability and Availability: Replication, caching, and distribution of directory data ensure high availability and fault tolerance, making the system robust against server failures.

4. Flexibility and Extensibility: The separation of directory trees and value trees allows the system to adapt to various types of resources beyond just IP addresses, supporting modern network needs.

5. Performance Optimization: Caching frequently requested names reduces latency and network traffic, speeding up name resolution and improving user experience.

6. Consistency and Management: Unique Directory Identifiers (DI) enable safe renaming, merging, and scalable management of names, avoiding conflicts in a global namespace.

7. Security Support: GNS supports secure naming through trusted directory management, access control mechanisms, and authenticated resource identifiers. These features help protect against unauthorized modifications, spoofing, and incorrect name resolution in distributed systems.

Chapter-5 Limitations of GNS

Despite its many strengths, the Global Name Service (GNS) has several inherent limitations and challenges:

1. Performance Overhead

Iterative Resolution: Name resolution in GNS involves iterative traversal across multiple directory levels and servers.

Latency Concerns: This process can lead to higher latency, especially when dealing with deep hierarchical structures or when authoritative servers are geographically distant.

Impact: The overhead affects response times compared to simpler flat namespaces, potentially impacting time-sensitive applications.

2. Management Complexity

Consistency Challenges: Managing directory merges, subtree moves, or zone transfers requires careful coordination.

Risk of Errors: Mistakes by administrators during these operations can introduce dangling references (broken links) or cycles within the directory tree.

Operational Overhead: These complexities increase the administrative burden and the risk of inconsistent or corrupted namespace states.

3. Vulnerability Risks

Trust Model: GNS relies on trusted authoritative servers to handle replication and caching.

Single Points of Failure: Compromise of primary servers can lead to propagation of stale or falsified data across replicas until the cache TTL (Time to Live) expires.

Security Threats: This makes the system vulnerable to attacks such as cache poisoning, spoofing, or denial-of-service, which can disrupt reliable name resolution.

Chapter-6 Real-World Implementation of GNS

The concepts underlying the Global Name Service (GNS) have found practical applications across various domains, particularly where scalable, secure, and decentralized name resolution is critical:

1. Academic and Experimental Systems

GNS principles have been implemented in numerous academic and experimental distributed systems. These platforms serve as testbeds to explore scalable and secure mechanisms for name resolution, helping to refine and advance the technology.

2. Decentralized Applications and Web3

In the rapidly growing blockchain and Web3 ecosystems, GNS is utilized to replace lengthy cryptographic addresses with user-friendly human-readable names. This improves usability and accessibility for decentralized services and digital identities.

3. Privacy-focused and Censorship-resistant Systems

GNS supports applications like GNUnet that prioritize privacy and resist centralized control. By avoiding reliance on centralized authorities, these systems provide greater freedom from censorship and enhanced user anonymity.

4. Internet of Things(IOT)

In IoT environments, GNS facilitates seamless identification of devices, sensors, and controllers even as they move between networks or change locations. This capability ensures persistent naming and reliable communication in dynamic, distributed IoT deployments.

Chapter-7 Conclusion

The Global Name Service (GNS) represents a significant advancement in naming and resource resolution for distributed systems. Its hierarchical and distributed architecture, supported by directory trees, value trees, unique Directory Identifiers, replication, and caching, ensures scalable and reliable name resolution.

GNS plays a critical role in modern computing environments such as cloud platforms, Internet of Things (IoT), and decentralized systems by enabling efficient service discovery, resource management, and secure naming. Despite its complexity and performance overhead, the flexibility and scalability of GNS make it a foundational component for next-generation distributed networks.

Reference

- [1] B. W. Lampson, M. Stiegler, and A. Finkelstein, "The Global Name Service: A distributed, scalable, and secure naming system," *ACM SIGOPS Operating Systems Review*, vol. 27, no. 1, pp. 1–12, Jan. 1993.
- [2] "DNS & Global Name Services," Scribd, [Online]. Available: <https://www.scribd.com/presentation/374144185/DNS-global-Name-Services1>. [Accessed: 22-Jan-2026].
- [3] R. Rana, "Introduction: Name Services and the Domain Name System, Directory and Discovery Services – Case Study of the Global Name Service," *Slideserve*, [Online]. Available: <https://www.slideserve.com/rane/introduction-name-services-and-the-domain-name-system-directory-and-discovery-services-case-study-of-the-global-name-ser>. [Accessed: 22-Jan-2026].
- [4] "Global Name Service (GNS)," Basic Computer Science, [Online]. Available: <https://basiccomputerscience.softicks.in/505/>. [Accessed: 22-Jan-2026].
- [5] "Global Name Service (GNS)," Brainkart, [Online]. Available: [https://www.brainkart.com/article/Global-Name-Service-\(GNS\)_8548/](https://www.brainkart.com/article/Global-Name-Service-(GNS)_8548/). [Accessed: 22-Jan-2026].